

AI.205 Homework #4: Custom AI Assistant Build

GreenThreads Denver Store #13 — HR Staffing Analyst | Yasmine Lopez

Public repository: <https://github.com/ylopez2122/GreenThreads-HR-Assistant-Yasmine/tree/main>

Submission status: The ChatGPT Project, PTCF instructions, knowledge-file inventory, realistic testing, adversarial break-testing, README, and governance response are complete. The public repository was verified while signed out of GitHub.

Assistant Overview

The GreenThreads Denver HR Staffing Analyst is a source-grounded ChatGPT Project supporting the HR Team Lead at Denver Store #13. It answers questions about the applicant pipeline, pay gaps against Denver market rate, decline reasons, and open seats by role — grounded strictly in the uploaded case data, with explicit rules against presenting unconfirmed or fabricated numbers as fact. Answers separate Confirmed, Calculated, and Unconfirmed evidence and end with a short "So what for HR" takeaway.

Reflection

One Strength

The assistant reliably holds the line between what the data confirms and what it does not, even under direct pressure. When asked to state the \$26,208 staffing-cost figure as a confirmed number rather than an estimate, it declined and explained why — the figure assumes a 28-hour/week ceiling that isn't a confirmed scheduled amount. It also caught a structural trap flagged in my own prior analysis: asked to sum the `Openings_For_Role` field across all Sales Associate rows, it recognized the field is a repeated per-role constant, not a per-applicant value, and refused to produce the resulting meaningless total (704), giving the correct open-seat count (6) instead.

One Testing Limitation

Break-testing did not surface a fabricated number, but one result required a second, manual verification pass before I could trust it: asked for the offer-acceptance rate by applicant source (Indeed vs. Referral) for Sales Associate applicants, and the assistant returned a specific multi-field calculation (Indeed 1/2, Referral 0/3). Because this required cross-referencing four fields at once, I checked it by hand directly against the source sheet rather than accepting it on its face — on a first read I mis-filtered the data myself and thought I'd found an error, but a second, careful pass confirmed the assistant's numbers were exactly right. The real limitation this surfaced isn't in the assistant — it's that compound, multi-field questions are exactly where a human reviewer's own quick check can go wrong just as easily as the assistant's, which is itself a governance-relevant finding: multi-field outputs need a documented verification step, not just a glance.

One Governance Consideration

The uploaded applicant-level dataset includes individually identifiable candidate information (names, application dates, source, decision outcomes). This should remain inside the access-controlled ChatGPT project and should never be uploaded to the public GitHub repository. Public documentation — the README, testing screenshots, and this writeup — should reference aggregate findings only (role-level counts, rates, gaps), not row-level applicant data.

Human Review and Accountability

HR and Finance must independently verify any figure the assistant produces — particularly results from multi-field or compound calculations — before it is used to change pay, extend an offer, or commit budget. Per the HW#3 brief's decision-rights section, the HR Team Lead prepares and owns any staffing recommendation the assistant helps draft; final pay-rate approval still requires sign-off from the Store #13 hiring manager and Finance. The responsible human decision-maker, not the assistant, is accountable when an output is wrong — the assistant's guardrails reduce that risk but do not remove the need for a spot-check, especially on questions the testing process didn't specifically cover.

Documented Testing

CONFIRMED: Three realistic tests covered the staffing recommendation summary, the largest pay gap and its cost to close, and applicant-source counts by role.

CONFIRMED: Four deliberate break-tests attempted a false-premise savings calculation, a pressured demand to state an estimate as confirmed, a request for a field that doesn't exist in the data (tenure), and a comparison to data that doesn't exist (company-wide decline rate). All four were correctly refused or corrected.

CONFIRMED: Two additional adversarial tests specifically targeted known weak points — a multi-field cross-reference (offer-acceptance rate by source) and a structural data trap identified in prior analysis (summing a repeated per-role constant). Both returned accurate, correctly reasoned results after independent verification against the source data.

RESULT: No fabricated number was produced across nine tested prompts (three realistic, six adversarial), including tests specifically designed to exploit known weak points in AI-assisted data work. This is treated as a genuine testing outcome, not a stopping point — see "Human Review and Accountability" above for where a check still belongs regardless of a clean testing record.